

Wine Reviews 数据集

该数据集有129971条数据，13个属性标签，分别为

country: 产出国

description: 描述

designation: 葡萄酒名称

points: 度数

price: 价格

province: 产出省

region_1: 产出区域1

region_2: 产出区域2

taster_name: 品鉴师

taster_twitter_handle: 品鉴师推特号

title: 品鉴师所获荣誉

variety: 品种

winery: 酒厂

```
In [61]: import os
import sys
import math
import pandas as pd
import numpy as np
import csv
import json
import pickle
```

```
import matplotlib.pyplot as plt
from scipy import stats
import statsmodels.api as sm
%matplotlib inline
```

```
In [4]: winel_data = pd.read_csv('..data\Wine Reviews\winemag-data-130k-v2.csv', index_col = 0)
winel_data.head()
```

```
Out[4]:
```

	country	description	designation	points	price	province	region_1	region_2	taster_name	taster_twitter_handle	title	variety	winery
0	Italy	Aromas include tropical fruit, broom, brimston...	Vulkà Bianco	87	NaN	Sicily & Sardinia	Etna	NaN	Kerin O'Keefe	@kerinokeefe	Nicosia 2013 Vulkà Bianco (Etna)	White Blend	Nicosia
1	Portugal	This is ripe and fruity, a wine that is smooth...	Avidagos	87	15.0	Douro	NaN	NaN	Roger Voss	@vossroger	Quinta dos Avidagos 2011 Avidagos Red (Douro)	Portuguese Red	Quinta dos Avidagos
2	US	Tart and snappy, the flavors of lime flesh and...	NaN	87	14.0	Oregon	Willamette Valley	Willamette Valley	Paul Gregutt	@paulgwine	Rainstorm 2013 Pinot Gris (Willamette Valley)	Pinot Gris	Rainstorm
3	US	Pineapple rind, lemon pith and orange blossom ...	Reserve Late Harvest	87	13.0	Michigan	Lake Michigan Shore	NaN	Alexander Peartree	NaN	St. Julian 2013 Reserve Late Harvest Riesling ...	Riesling	St. Julian
4	US	Much like the regular bottling from 2012, this...	Vintner's Reserve Wild Child Block	87	65.0	Oregon	Willamette Valley	Willamette Valley	Paul Gregutt	@paulgwine	Sweet Cheeks 2012 Vintner's Reserve Wild Child...	Pinot Noir	Sweet Cheeks

```
In [15]: wine1_data.shape
```

```
Out[15]: (129971, 13)
```

```
In [14]: cols = list(wine1_data)
        cols
```

```
Out[14]: ['country',
          'description',
          'designation',
          'points',
          'price',
          'province',
          'region_1',
          'region_2',
          'taster_name',
          'taster_twitter_handle',
          'title',
          'variety',
          'winery']
```

数据摘要

对标称数据计算频数

标称属性包括: 'country', 'designation', 'province', 'region_1', 'region_2', 'taster_name', 'taster_twitter_handle', 'variety', 'winery', 分别计算它们的频数

```
In [25]: nominal_attribute = ['country', 'designation', 'province', 'region_1', 'region_2', 'taster_name', 'taster_twitter_handle', 'variety', 'winery']
        for tmp in nominal_attribute:
            print(wine1_data[tmp].value_counts())
            print('-' * 60)
```

US	54504
France	22093
Italy	19540
Spain	6645
Portugal	5691
Chile	4472
Argentina	3800
Austria	3345
Australia	2329
Germany	2165
New Zealand	1419

South Africa	1401
Israel	505
Greece	466
Canada	257
Hungary	146
Bulgaria	141
Romania	120
Uruguay	109
Turkey	90
Slovenia	87
Georgia	86
England	74
Croatia	73
Mexico	70
Moldova	59
Brazil	52
Lebanon	35
Morocco	28
Peru	16
Ukraine	14
Macedonia	12
Serbia	12
Czech Republic	12
Cyprus	11
India	9
Switzerland	7
Luxembourg	6
Bosnia and Herzegovina	2
Armenia	2
Egypt	1
Slovakia	1
China	1

Name: country, dtype: int64

Reserve	2009
Estate	1322
Reserva	1259
Riserva	698
Estate Grown	621

...

Private Stash #10	1
Ten Degrees Vineyard	1
CLB Reserve	1
Million Dollar Beach	1
Faiv Brut Rosé Metodo Classico	1

Name: designation, Length: 37979, dtype: int64

California	36247
Washington	8639
Bordeaux	5941
Tuscany	5897
Oregon	5373

...

Pitsilia Mountains	1
Markopoulo	1
Elazığ-Diyarbakir	1
China	1
Krk	1

Name: province, Length: 425, dtype: int64

Napa Valley	4480
Columbia Valley (WA)	4124
Russian River Valley	3091
California	2629
Paso Robles	2350

...

Cabernet de Saumur	1
Riverland	1
Mâcon-Mancey	1
Jujuy	1
Gippsland	1

Name: region_1, Length: 1229, dtype: int64

Central Coast	11065
Sonoma	9028
Columbia Valley	8103
Napa	6814
Willamette Valley	3423
California Other	2663
Finger Lakes	1777
Sierra Foothills	1462
Napa-Sonoma	1169
Central Valley	1062
Southern Oregon	917
Oregon Other	727
Long Island	680
North Coast	584
Washington Other	534
South Coast	272
New York Other	231

Name: region_2, dtype: int64

Roger Voss	25514
Michael Schachner	15134

Kerin O' Keefe	10776
Virginie Boone	9537
Paul Gregutt	9532
Matt Kettmann	6332
Joe Czerwinski	5147
Sean P. Sullivan	4966
Anna Lee C. Iijima	4415
Jim Gordon	4177
Anne Krebiehl MW	3685
Lauren Buzzeo	1835
Susan Kostrzewa	1085
Mike DeSimone	514
Jeff Jenssen	491
Alexander Peartree	415
Carrie Dykes	139
Fiona Adams	27
Christina Pickard	6

Name: taster_name, dtype: int64

@vossroger	25514
@wineschach	15134
@kerinokeefe	10776
@vboone	9537
@paulgwine	9532
@mattkettmann	6332
@JoeCz	5147
@wawinereport	4966
@gordone_cellars	4177
@AnneInVino	3685
@laurbuzz	1835
@suskostrzewa	1085
@worldwineguys	1005
@bkfiona	27
@winewchristina	6

Name: taster_twitter_handle, dtype: int64

Pinot Noir	13272
Chardonnay	11753
Cabernet Sauvignon	9472
Red Blend	8946
Bordeaux-style Red Blend	6915
...	
Moschofilero-Chardonnay	1
Francisa	1
Sercial	1
Ondenc	1
Merlot-Argaman	1

```
Name: variety, Length: 707, dtype: int64
```

```
-----
Wines & Winemakers      222
Testarossa              218
DFJ Vinhos              215
Williams Selyem         211
Louis Latour            199
```

```
...
Château La Croix Lartigue 1
Geode                    1
Patrick M. Paul          1
Heredad Soliterra        1
Once Upon a Vine         1
```

```
Name: winery, Length: 16757, dtype: int64
```

- 从数据中发现，美国为最大产出国
- California是最大的产出州
- 名字为Roger Voss的品鉴师，品鉴次数最多，他的推特号为@vossroger
- Pinot Noir种类最多

对数值数据计算五数概括以及缺失值

```
In [27]: number_data = ['points', 'price']
         wine1_data[number_data].describe()
```

```
Out[27]:
```

	points	price
count	129971.000000	120975.000000
mean	88.447138	35.363389
std	3.039730	41.022218
min	80.000000	4.000000
25%	86.000000	17.000000
50%	88.000000	25.000000
75%	91.000000	42.000000
max	100.000000	3300.000000

```
In [41]: wine1_data.isnull()[number_data].sum()
```


```
Out[41]: points      0  
price      8996  
dtype: int64
```


数值数据包括'point' 和 'price'

- point: 最大100、最小80、Q1值86、中位数88、Q3值91, 缺失值个数为0
- price: 最大3300、最小4、Q1值17、中位数25、Q3值42, 缺失值个数为0


数据可视化

```
In [10... wine1_data['points'].hist()  
  
points = wine1_data['points'].dropna()  
points = points.apply(lambda x: x + np.random.normal())  
  
fig = plt.figure()  
res = stats.probplot(points, plot=plt)  
plt.show()  
  
wine1_data.boxplot(column=['points'])
```





Out[105]: <matplotlib.axes._subplots.AxesSubplot at 0x25823452048>



我们可以发现'points'数据符合正态分布

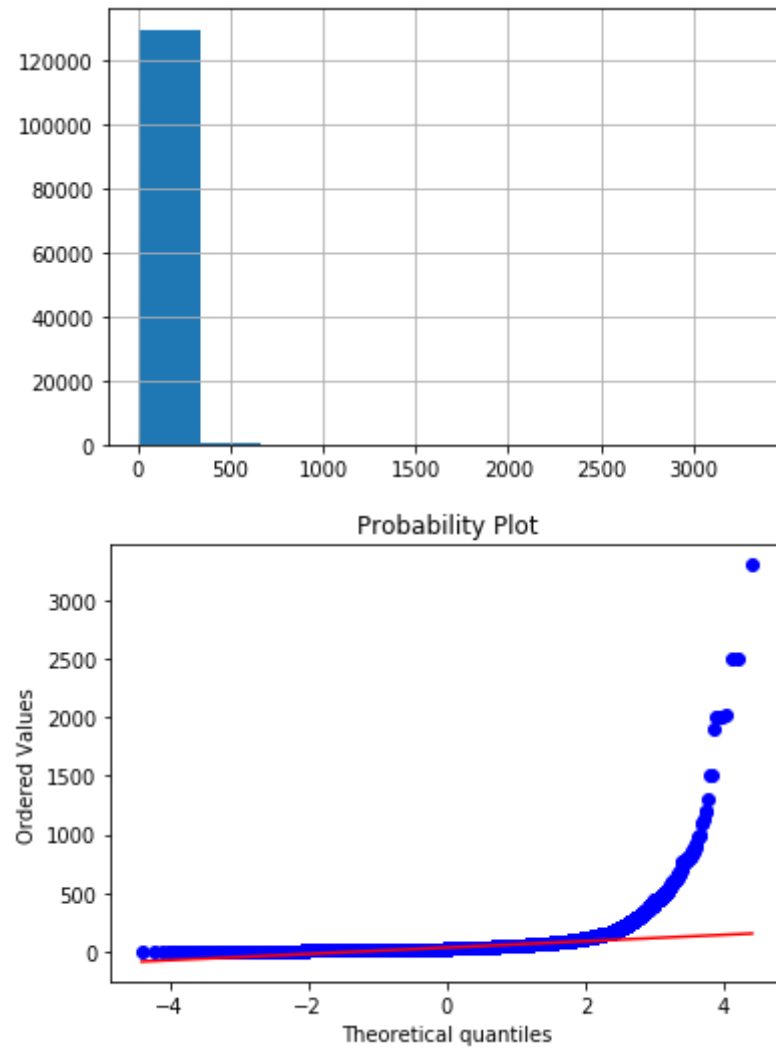
```
In [10... wine1_data['price'].hist()

price = wine1_data['price'].dropna()
price = price.apply(lambda x: x + np.random.normal())


fig = plt.figure()
```

```
res = stats.probplot(price, plot=plt)
plt.show()

winel_data.boxplot(column=['price'])
```



```
Out[106]: <matplotlib.axes._subplots.AxesSubplot at 0x258234f9188>
```



'price' 数据符合正态分布，高价酒较少，价格主要集中在中低价位

数据缺失的处理

```
In [43]: wine1_data.isnull()[cols].sum()
```

```
Out[43]: country          63
description          0
designation        37465
points              0
price              8996
province            63
region_1           21247
region_2           79460
taster_name        26244
taster_twitter_handle 31213
title               0
variety              1
winery              0
dtype: int64
```


- 对于country和province的缺失，可能无法确定该葡萄酒的产出国
- taster_name缺失可能说明该葡萄酒没有品酒师去品鉴
- taster_twitter_handle缺失说明品酒师没有获得任何荣誉称号

将缺失部分剔除

```
In [44]: delete_wine1 = wine1_data.dropna()
```


```
In [47]: delete_wine1['points'].hist()
```

```
Out[47]: <matplotlib.axes._subplots.AxesSubplot at 0x2581ae56208>
```




```
In [48]: wine1_data['points'].hist()
```

```
Out[48]: <matplotlib.axes._subplots.AxesSubplot at 0x2581ae95388>
```




```
In [49]: delete_wine1['price'].hist()
```

```
Out[49]: <matplotlib.axes._subplots.AxesSubplot at 0x2581af9f048>
```



```
In [50]: wine1_data['price'].hist()
```

```
Out[50]: <matplotlib.axes._subplots.AxesSubplot at 0x2581afc9188>
```




剔除缺失值后，prices的分布没有太多改变

用最高频率值来填补缺失值

```
In [52]: fill_max = wine1_data.fillna({'points': wine1_data['points'].mode().item(), 'price': wine1_data['price'].mode().item()})
```


```
In [53]: fill_max['points'].hist()
```

```
Out[53]: <matplotlib.axes._subplots.AxesSubplot at 0x2581bc8f448>
```



```
In [54]: fill_max['price'].hist()
```

```
Out[54]: <matplotlib.axes._subplots.AxesSubplot at 0x2581bd85948>
```



通过属性的相关关系来填补缺失值

首先计算属性之间的相关系数


```
In [59]: x = winel_data.corr()  
print(x)
```

```
           points    price  
points  1.000000  0.416167  
price   0.416167  1.000000
```

只有price和points为数值属性，虽然相关性为中，但是points没有缺失数据，可以根据已有的points和price数据，得到它们的回归方程，利用回归方程计算缺失值

```
In [10... # 进行回归  
points = winel_data['points']  
price = winel_data['price']  
model = sm.OLS(price, points).fit()  
new_data = winel_data  
for i in range(len(new_data)):  
    if np.isnan(new_data['price'][i]):  
        new_data.loc[i, 'price'] = model.predict(new_data['points'][i])  
number_data = ['points', 'price']  
new_data.hist()
```

```
Out[108]: array([[<matplotlib.axes._subplots.AxesSubplot object at 0x00000258234C2888>,  
                  <matplotlib.axes._subplots.AxesSubplot object at 0x0000025823CD6F48>]],  
              dtype=object)
```



```
In [10... new_data[number_data].describe()
```

Out[109]:

	points	price
count	129971.000000	129970.000000
mean	88.447138	35.447080
std	3.039730	41.080914
min	80.000000	4.000000
25%	86.000000	17.000000
50%	88.000000	25.000000
75%	91.000000	42.000000
max	100.000000	3300.000000

填充后，price只有均值和标准差发生变化，均值减小。

通过数据对象之间的相似性来填补缺失值

根据对象之间ponits的相似性，填充缺失的price


```
In [ ]: df_sim = wine1_data[['price', 'points']]
p = {}
for row in df_sim.iterrows():
    if p.get(row[1]['points'], None):
        if not np.isnan(row[1]['price']):
            p[row[1]['points']][0] += row[1]['price']
            p[row[1]['points']][1] += 1
    else:
        if not np.isnan(row[1]['price']):
            p[row[1]['points']] = [row[1]['price'], 1]
for k in p.keys():
    p[k][0] = round(p[k][0] / p[k][1], 4)
for i in range(len(df_sim['price'])):
    if (np.isnan(df_sim['price'][i])):
        da = p[df_sim.loc[i, 'points']][0]
        df_sim.loc[i, 'price'] = da
```

```
In [10... number_data = ['points', 'price']
df_sim[number_data].describe()
```

```
Out[104]:
```

	points	price
count	129971.000000	129971.000000
mean	88.447138	35.446999
std	3.039730	41.080766
min	80.000000	4.000000
25%	86.000000	17.000000
50%	88.000000	25.000000
75%	91.000000	42.000000
max	100.000000	3300.000000

填充后，price只有均值和标准差发生变化，均值增大。